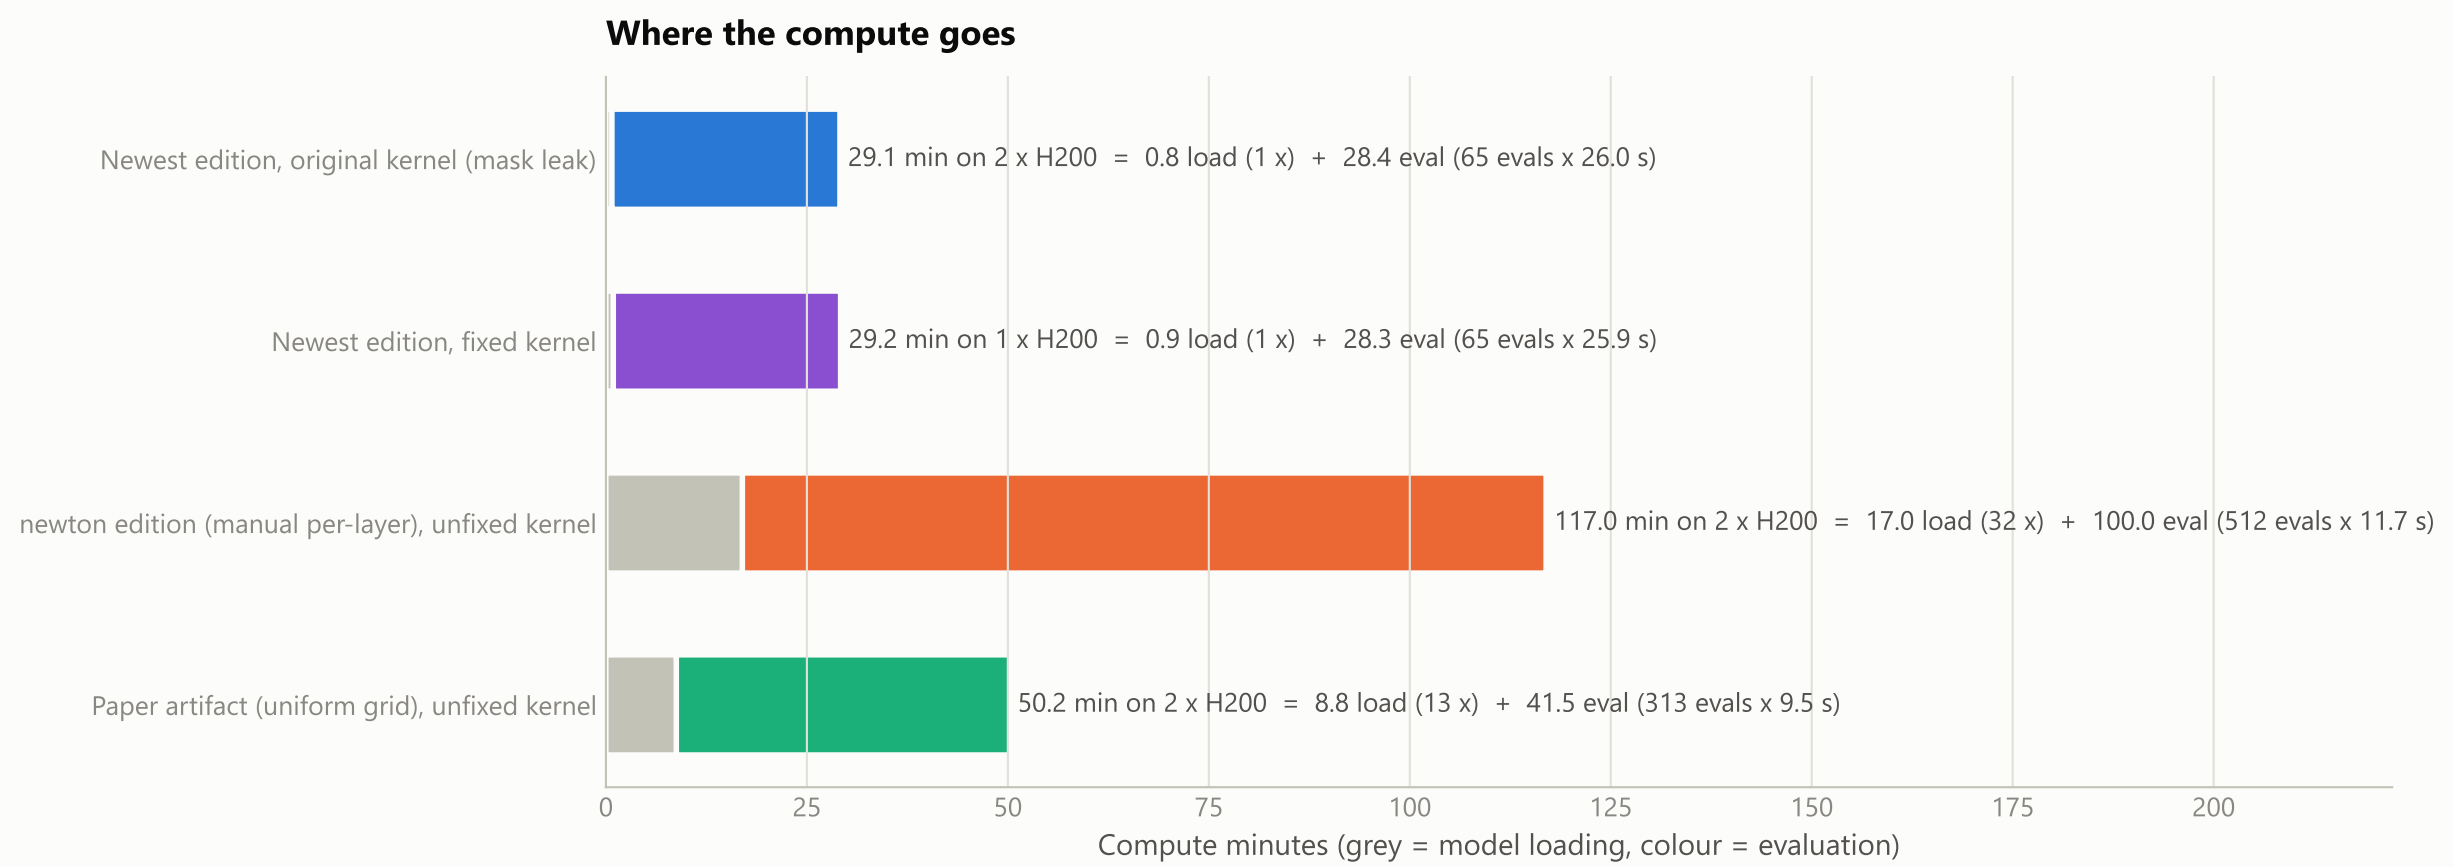
Where the compute goes
Newest edition, original kernel (mask leak)
29.1 min on 2 x H200 = 0.8 load (1 x) + 28.4 eval (65 evals x 26.0 s)
Newest edition, fixed kernel
29.2 min on 1 x H200 = 0.9 load (1 x) + 28.3 eval (65 evals x 25.9 s)
newton edition (manual per-layer), unfixed kernel
117.0 min on 2 x H200 = 17.0 load (32 x) + 100.0 eval (512 evals x 11.7 s)
Paper artifact (uniform grid), unfixed kernel
50.2 min on 2 x H200 = 8.8 load (13 x) + 41.5 eval (313 evals x 9.5 s)
0
25
50
75
100
125
150
175
200
Compute minutes (grey = model loading, colour = evaluation)