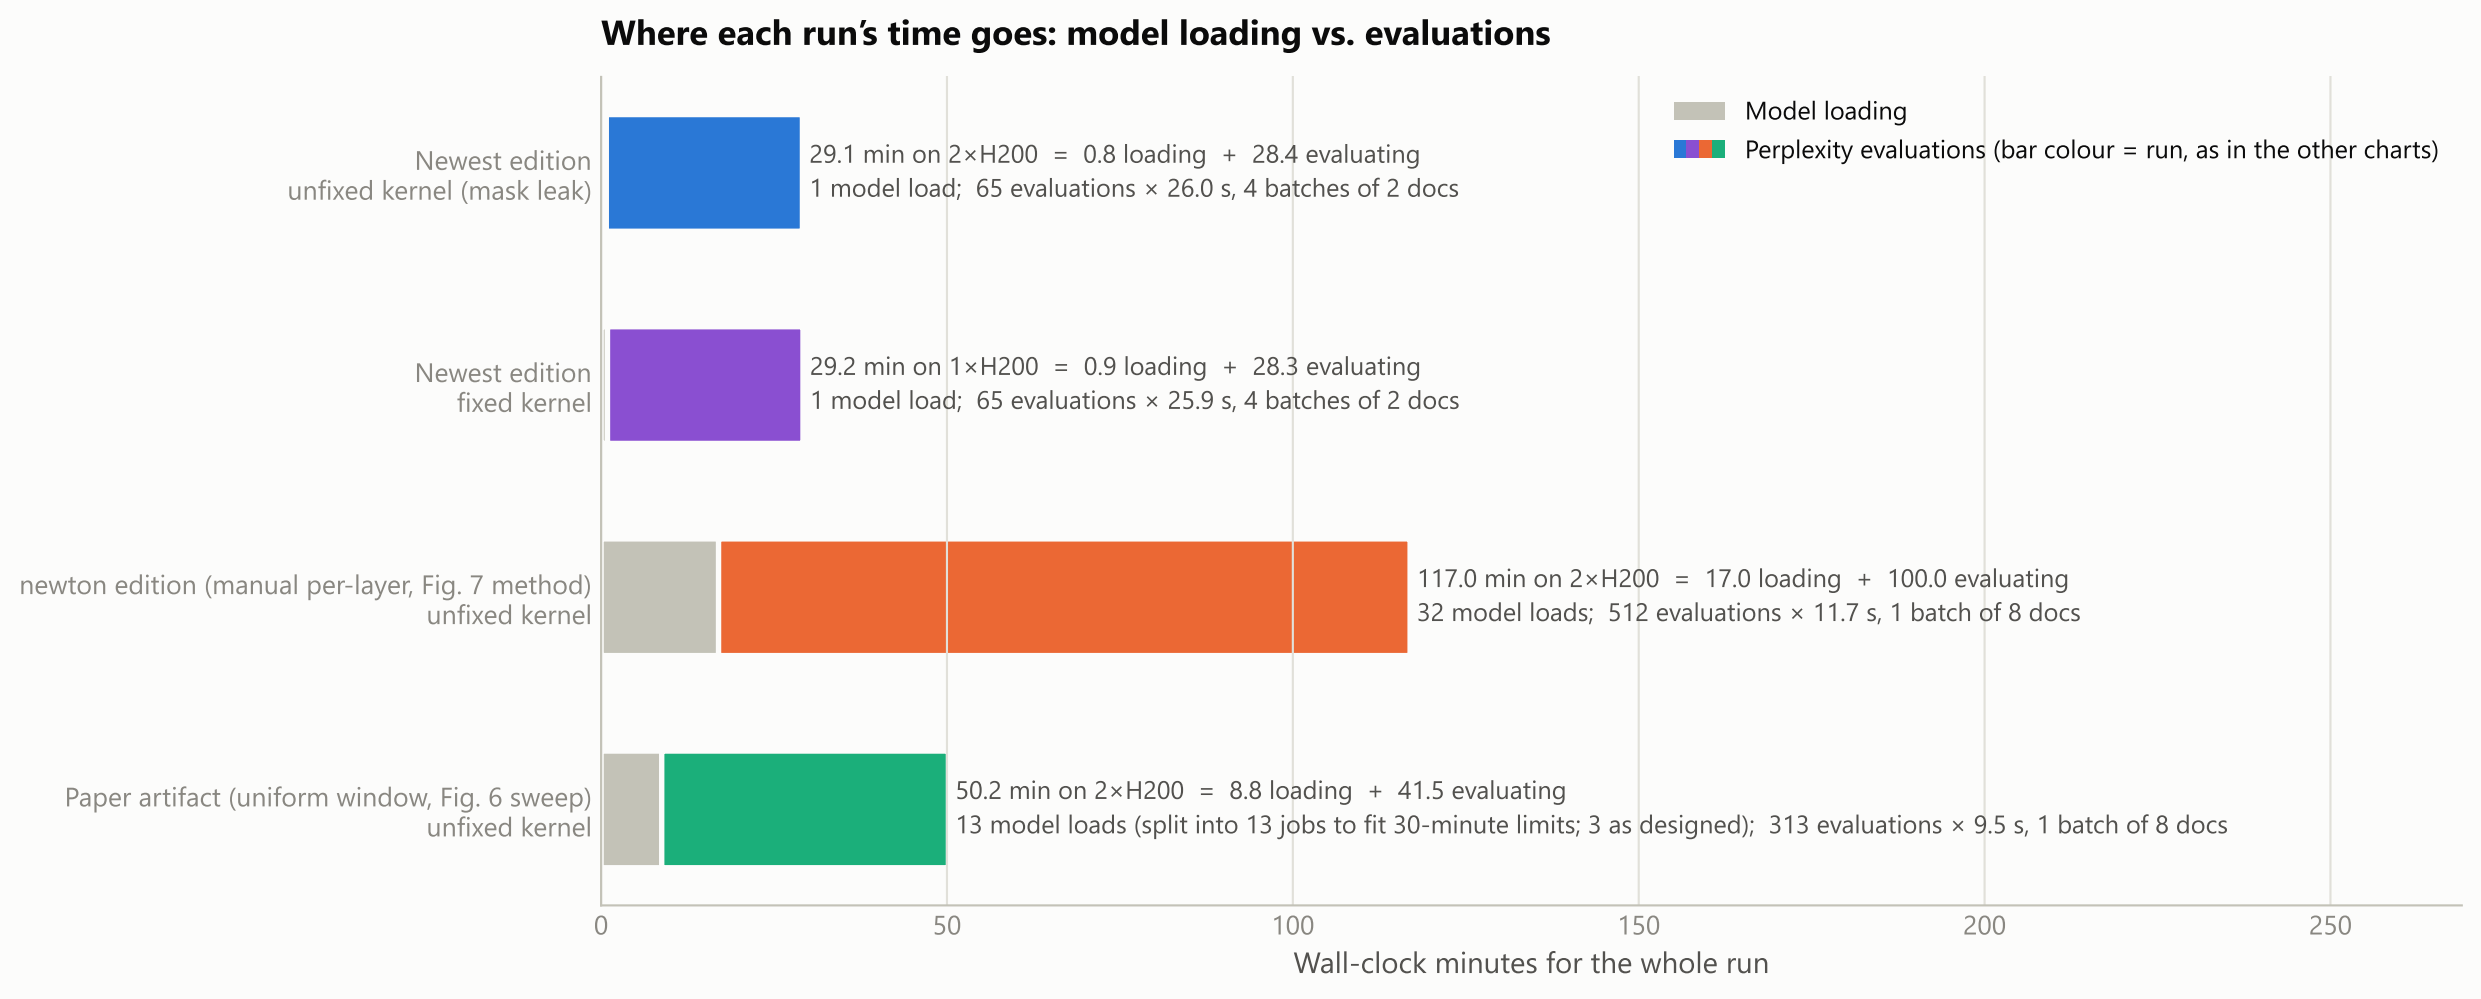
Where each run's time goes: model loading vs. evaluations
Model loading
Perplexity evaluations (bar colour = run, as in the other charts)
Newest edition
unfixed kernel (mask leak)
29.1 min on 2×H200 = 0.8 loading + 28.4 evaluating
1 model load; 65 evaluations × 26.0 s, 4 batches of 2 docs
Newest edition
fixed kernel
29.2 min on 1×H200 = 0.9 loading + 28.3 evaluating
1 model load; 65 evaluations × 25.9 s, 4 batches of 2 docs
newton edition (manual per-layer, Fig. 7 method)
unfixed kernel
117.0 min on 2×H200 = 17.0 loading + 100.0 evaluating
32 model loads; 512 evaluations × 11.7 s, 1 batch of 8 docs
Paper artifact (uniform window, Fig. 6 sweep)
unfixed kernel
50.2 min on 2×H200 = 8.8 loading + 41.5 evaluating
13 model loads (split into 13 jobs to fit 30-minute limits; 3 as designed); 313 evaluations × 9.5 s, 1 batch of 8 docs
0
50
100
150
200
250
Wall-clock minutes for the whole run