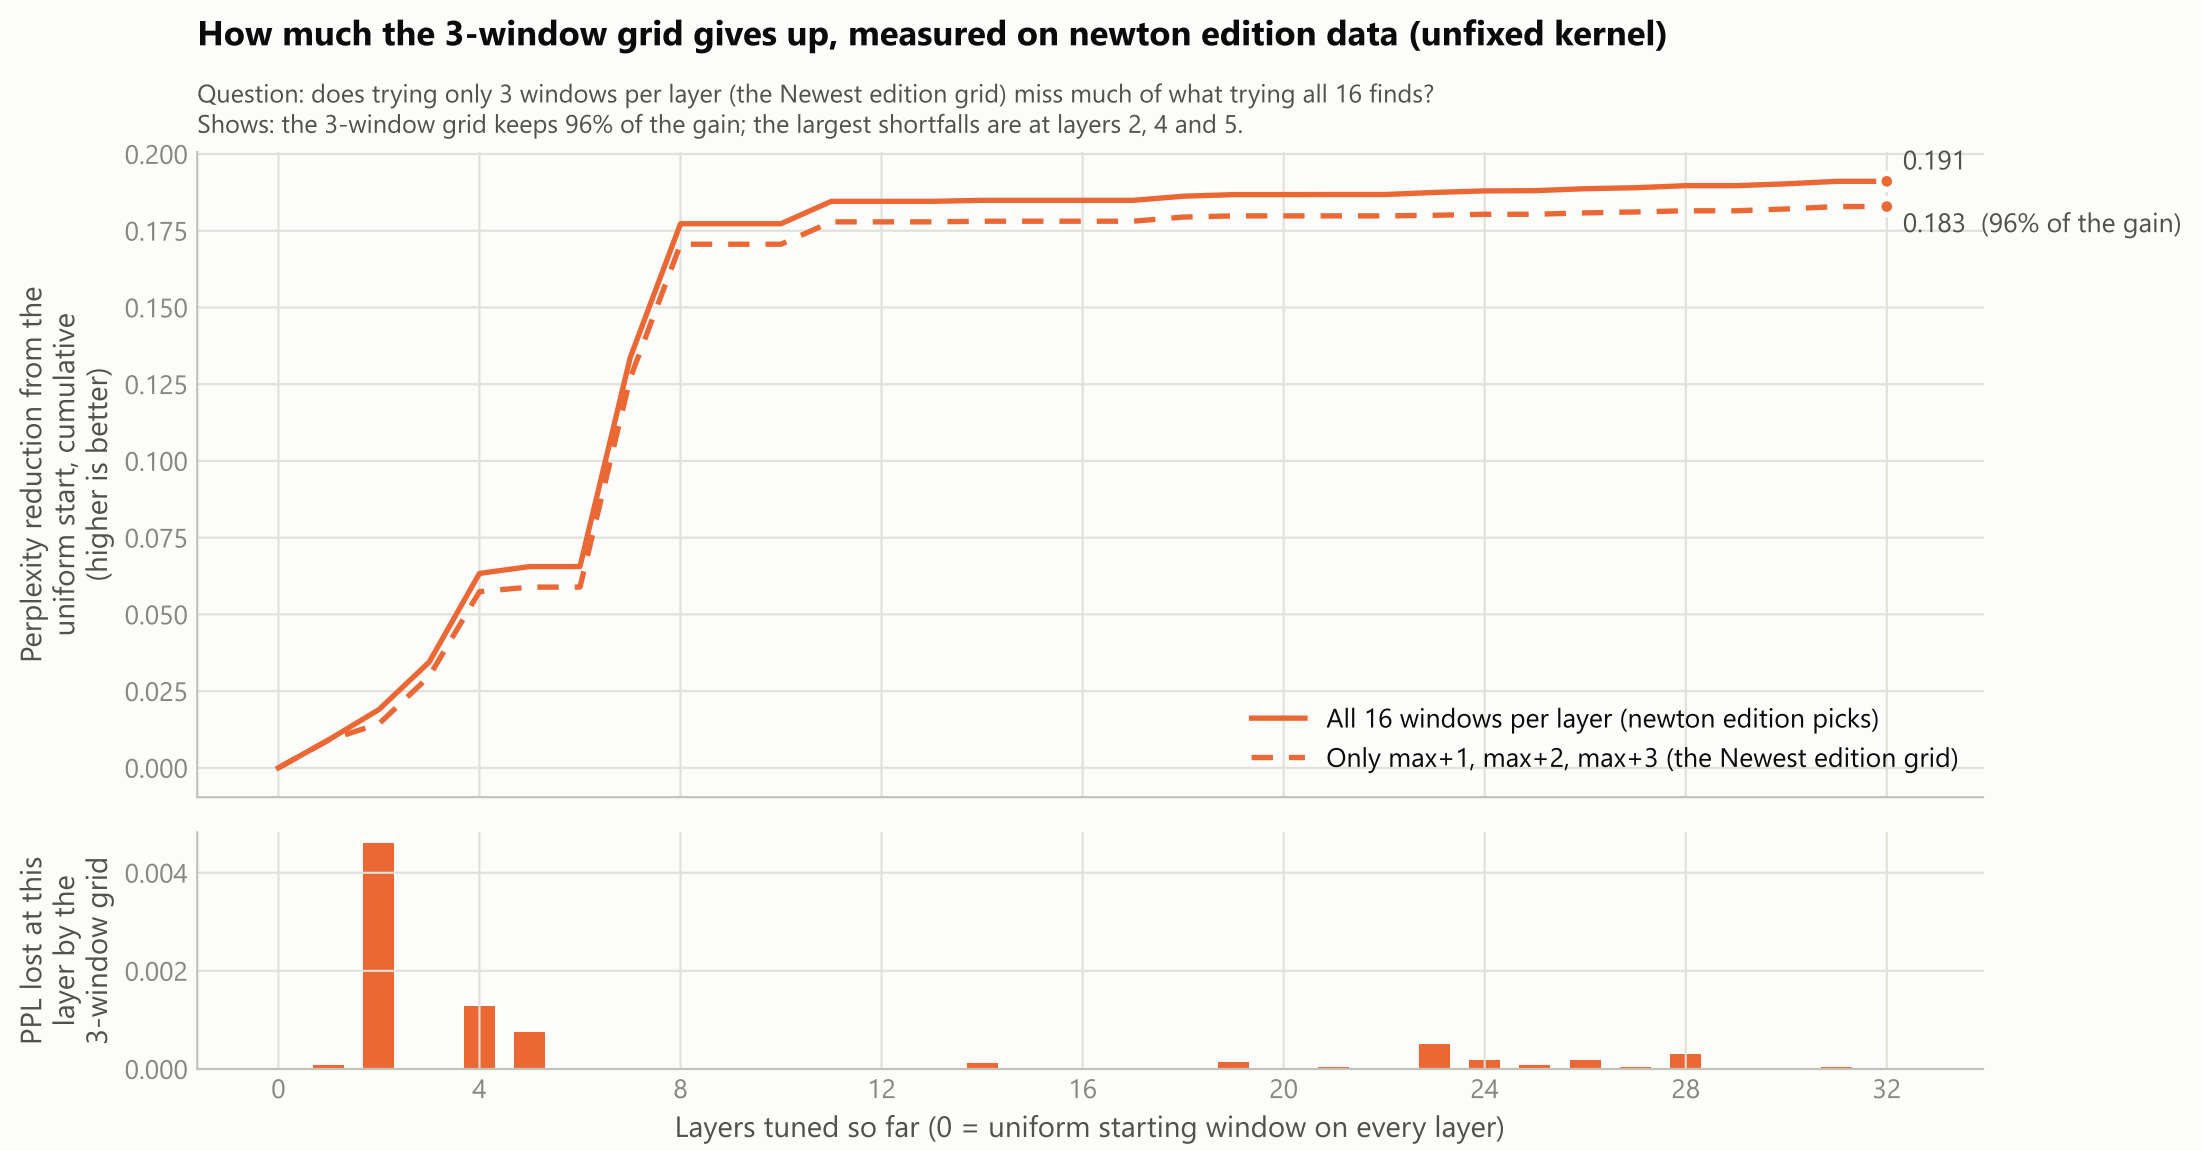
How much the 3-window grid gives up, measured on newton edition data (unfixed kernel)
Question: does trying only 3 windows per layer (the Newest edition grid) miss much of what trying all 16 finds?
Shows: the 3-window grid keeps 96% of the gain; the largest shortfalls are at layers 2, 4 and 5.
Perplexity reduction from the uniform start, cumulative (higher is better)
0.200
0.175
0.150
0.125
0.100
0.075
0.050
0.025
0.000
0.191
0.183 (96% of the gain)
All 16 windows per layer (newton edition picks)
Only max+1, max+2, max+3 (the Newest edition grid)
PPL lost at this layer by the 3-window grid
0.004
0.002
0.000
0
4
8
12
16
20
24
28
32
Layers tuned so far (0 = uniform starting window on every layer)
Approximate: each layer's windows were scored with the earlier layers at the newton edition's own picks.