Unfixed kernel (mask leak): both runs stay 0.350 to 0.835 PPL above the unpatched model

Question: how much perplexity does choosing a window per layer recover, and at which layers?
Each line is one pass over all 32 layers; each drop is a layer whose window the search changed.

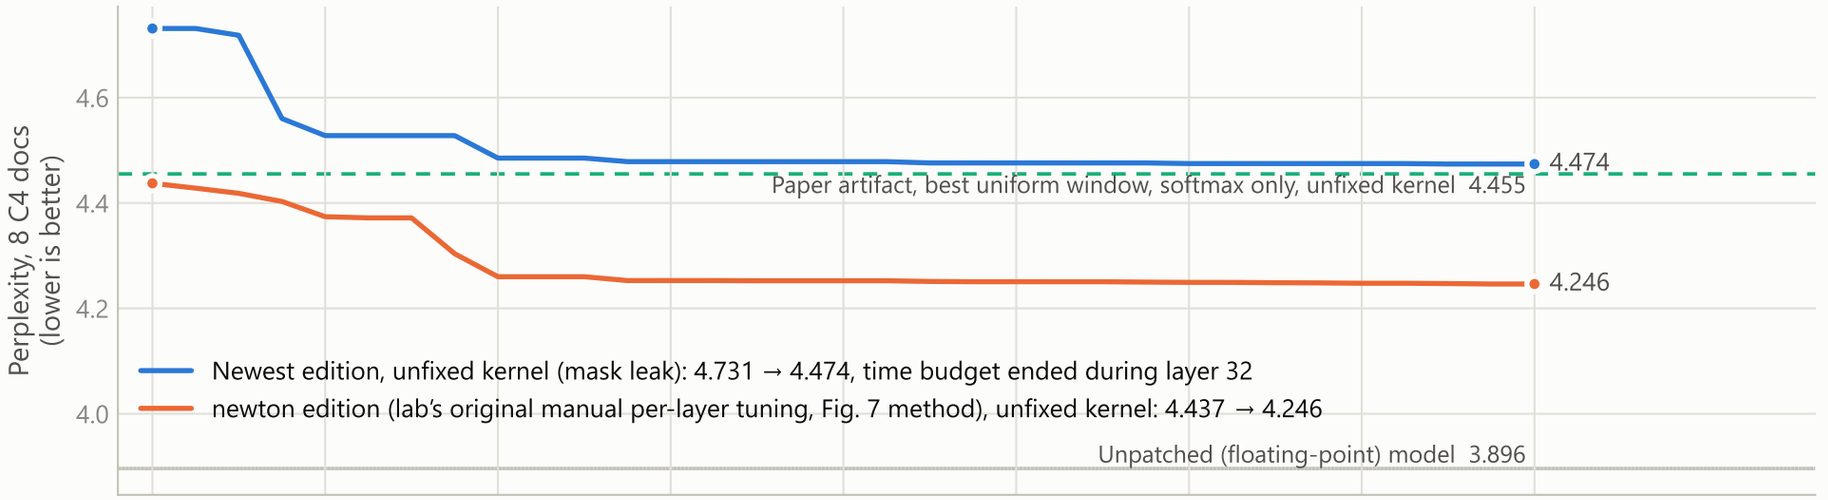

Fixed kernel, zoomed in: starts 0.0148 PPL above the unpatched model; search removes 0.0021

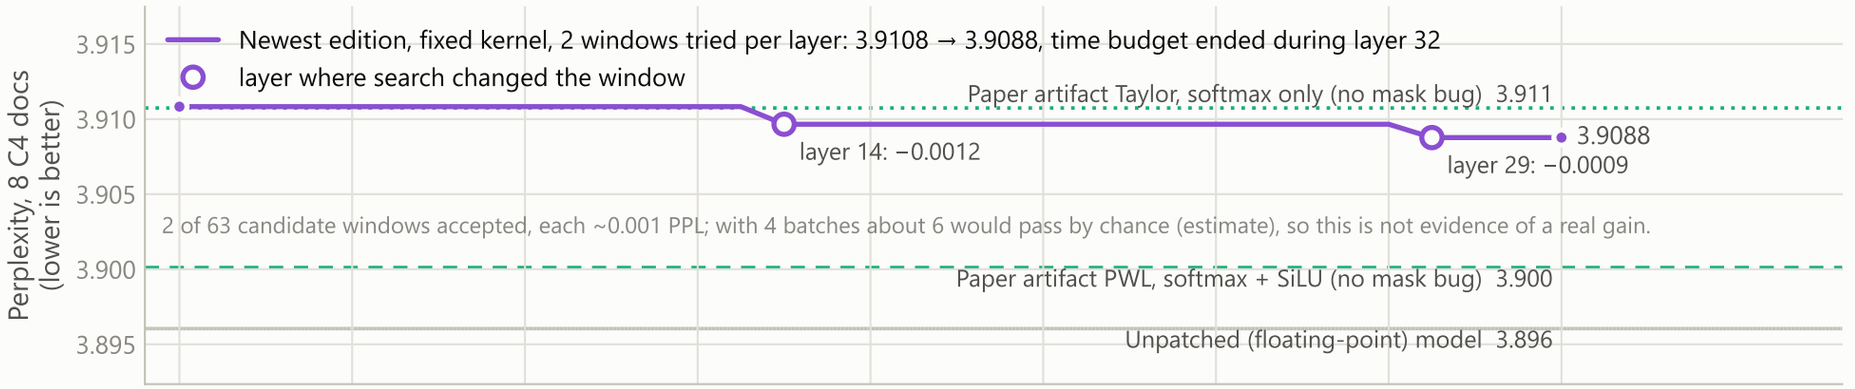

All runs on one scale: share of each run's own gap to the unpatched model that search closed
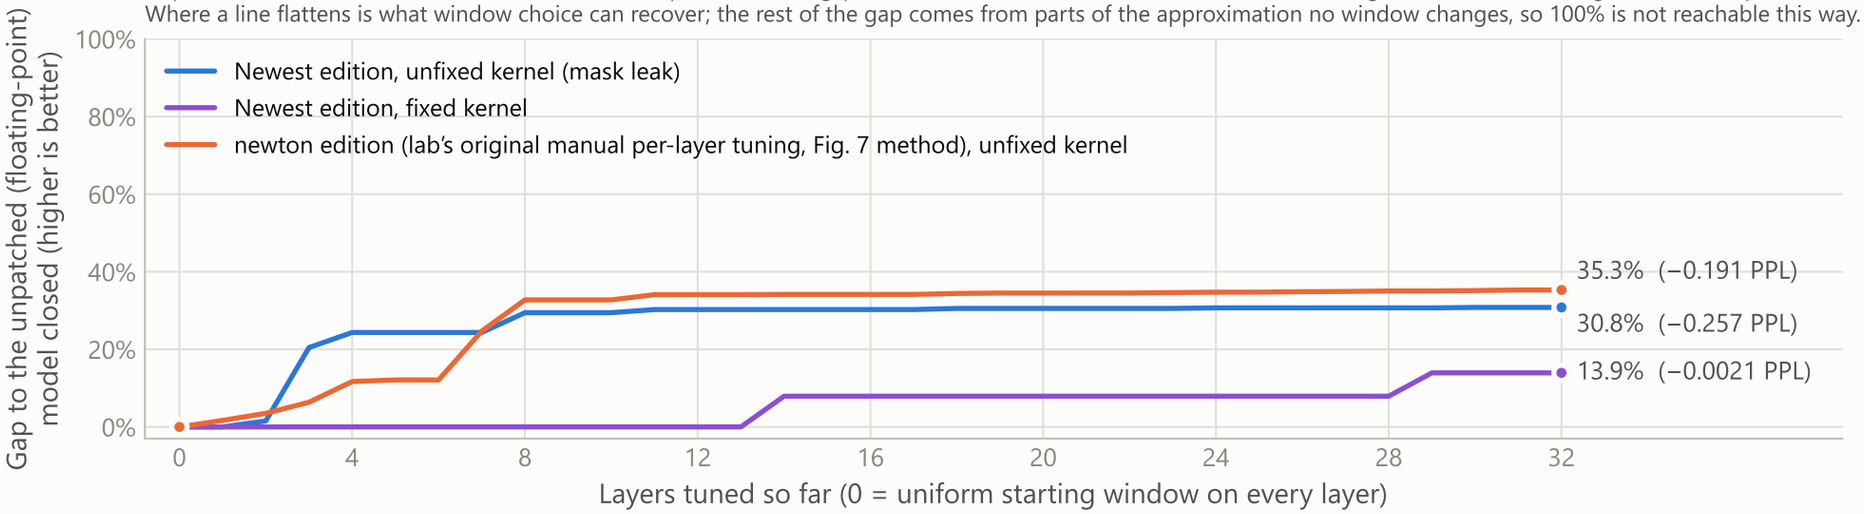
Gap closed = (start PPL − PPL now) / (start PPL − unpatched floating-point PPL). 0% = no better than the uniform starting window, 100% = as good as the unpatched model.
Where a line flattens is what window choice can recover; the rest of the gap comes from parts of the approximation no window changes, so 100% is not reachable this way.